



Long-horizon training: per-step peak memory and whole-step time (running average)


0.84B GPU | retained graph, resident, in-backward (fastest) — memory




0.84B GPU | retained graph, resident, in-backward (fastest) — time




0.84B GPU | recomputation, resident, deferred — memory




0.84B GPU | recomputation, resident, deferred — time




0.84B GPU | recomputation, streamed, deferred (memory floor) — memory




0.84B GPU | recomputation, streamed, deferred (memory floor) — time




0.84B GPU | recomputation, streamed, in-backward — memory




0.84B GPU | recomputation, streamed, in-backward — time




6.9B GPU | recomputation, resident, deferred — memory
+2.92e4




6.9B GPU | recomputation, resident, deferred — time




6.9B GPU | recomputation, streamed, in-backward — memory




6.9B GPU | recomputation, streamed, in-backward — time




0.84B CPU | recomputation, streamed, deferred — memory




0.84B CPU | recomputation, streamed, deferred — time




0.84B CPU | retained graph, resident, in-backward — memory



0.84B CPU | retained graph, resident, in-backward — time



6.9B CPU | recomputation, streamed, in-backward — memory



6.9B CPU | recomputation, streamed, in-backward — time

