



Long-horizon serving: per-request peak memory and latency (running average)


0.84B GPU serving | resident, KV cache — memory




0.84B GPU serving | resident, KV cache — time




0.84B GPU serving | streamed, recompute — memory




0.84B GPU serving | streamed, recompute — time




0.84B CPU serving | ONNX runtime, disk-streamed weights — memory



0.84B CPU serving | ONNX runtime, disk-streamed weights — time



6.9B GPU serving | resident, KV cache — memory



6.9B GPU serving | resident, KV cache — time

