

Segmented execution: one segment executing at a time



GPT decoder — four segmentation axes

E — embedding
split d_{model} , concat

A — attention
head groups, shared W_o

M — MLP
 d_{ff} chunks, running-sum

H — output head
vocab slices, streamed CE




evict



load



park



dials: recomputation · weight streaming · update schedule · dropout · KV cache $\Rightarrow$ 12 training + 4 inference modes

every mode computes the same model — identical weights, only the memory schedule changes